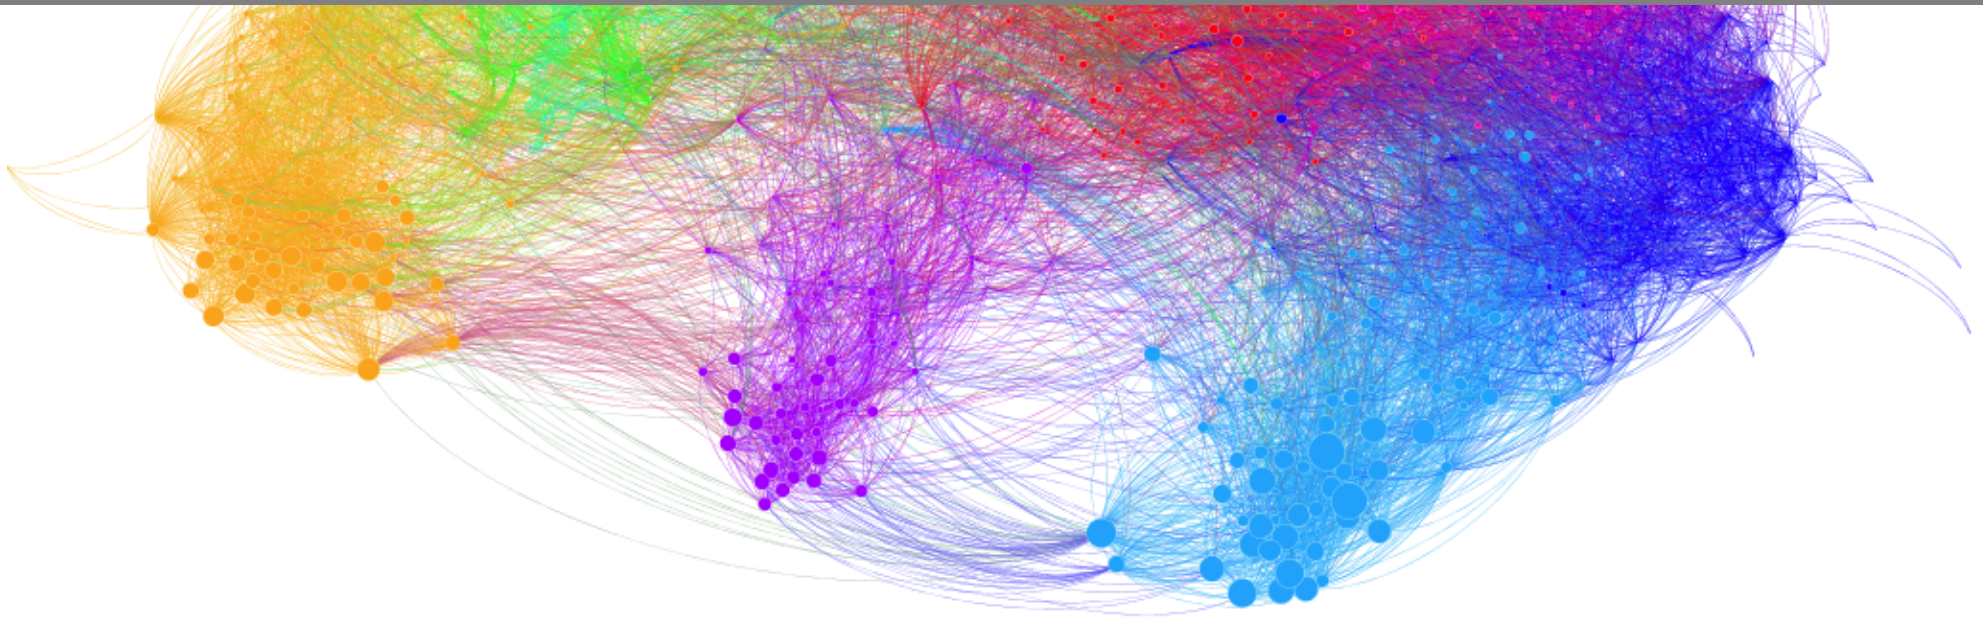


Engineering High-Performance Community Detection Heuristics for Massive Graphs


Christian L. Staudt and Henning Meyerhenke

INSTITUTE OF THEORETICAL INFORMATICS · PARALLEL COMPUTING GROUP



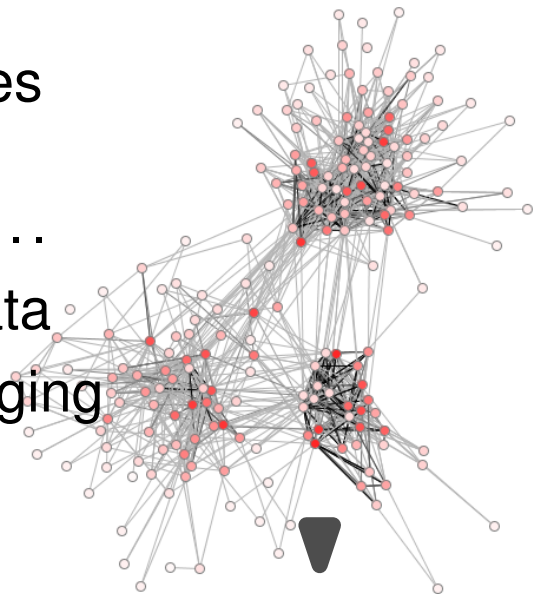
Big Network Data

- proliferation of large networks and high data rates
 - e.g. WWW (> 30 billion pages),
online social networks (> 600 million active users),...
- analysts need information from these piles of data
- **complex networks** are computationally challenging
 - **scale-free** topology \rightarrow load balancing issues
 - **small-world** network \rightarrow cache performance issues



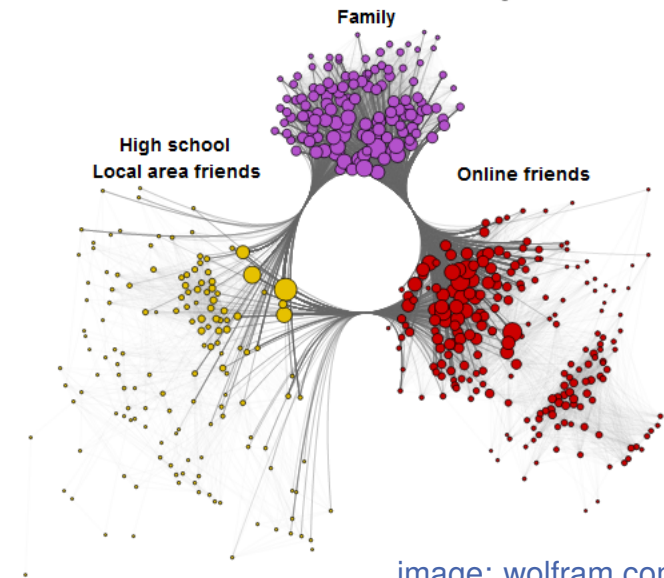
Big Network Data

- proliferation of large networks and high data rates
 - e.g. WWW (> 30 billion pages),
online social networks (> 600 million active users),...
- analysts need information from these piles of data
- **complex networks** are computationally challenging
 - **scale-free** topology \rightarrow load balancing issues
 - **small-world** network \rightarrow cache performance issues



Community Detection

- find **internally dense, externally sparse subgraphs** (formalized: e.g. **modularity**)
- goals: uncover community structure, prepartition network




[survey: Schaeffer 07, Fortunato 10]

image: wolfram.com

Challenge


10th DIMACS Implementation Challenge - Graph Partitioning and Clustering



- criteria: time and quality (**modularity**)
- high-quality solutions
 - **RG**, a randomized greedy agglomerative algorithm
 - **CGGC**, an ensemble using **RG** [Ovelgönne & Geyer-Schulz 13]
- large variance in running time among contestants
- few relied on parallelism
 - **CLU_TBB**, a parallel agglomerative algorithm [Fagginger Auer, Bisseling 13]
- few could handle largest graphs (billions of edges)

Challenge

10th DIMACS Implementation Challenge *- Graph Partitioning and Clustering*



- criteria: time and quality (**modularity**)
- high-quality solutions
 - **RG**, a randomized greedy agglomerative algorithm
 - **CGGC**, an ensemble using **RG** [Ovelgönne & Geyer-Schulz 13]
- large variance in running time among contestants
- few relied on parallelism
 - **CLU_TBB**, a parallel agglomerative algorithm [Fagginger Auer, Bisseling 13]
- few could handle largest graphs (billions of edges)

Others


- original label propagation algorithm [Raghavan et al. 07]
- distributed parallel label propagation on Hadoop [Ovelgönne 12]

Requirements

- only nearly **linear time** algorithms are practical
- we need to take advantage of **parallelism**

Our Approach

- algorithm engineering
- a framework of **shared-memory parallel heuristics**
 - **PLP**: a **label propagation** algorithm
 - **PLM**: a parallelization of a locally greedy **modularity maximizer**
 - **PLMR**: **PLM** with additional refinement
 - **EPP**: an **ensemble technique**




Requirements

- only nearly **linear time** algorithms are practical
- we need to take advantage of **parallelism**

Our Approach

- algorithm engineering
- a framework of **shared-memory parallel heuristics**
 - **PLP**: a **label propagation** algorithm
 - **PLM**: a parallelization of a locally greedy **modularity maximizer**
 - **PLMR**: **PLM** with additional refinement
 - **EPP**: an **ensemble technique**



Capabilities

- **PLM** first parallel variant of **Louvain method**, optional refinement phase
- data rates approach 50M edges/sec, depending on algorithm
- **NetworKit**: a framework for high-performance network analysis


- Objective function **modularity**:

$$q(\mathcal{C}) = \sum_{C \in \mathcal{C}} \left(\frac{|E(C)|}{m} - \left(\frac{\sum_{v \in C} \deg(v)}{2m} \right)^2 \right)$$


- Expected deviation from random graph with the same degree sequence
- **NP-hard** to optimize for modularity [Brandes et al., IEEE TKDE 2008] and most other (interesting) objective functions
- Modularity has some **known issues** (resolution limit, ...), some can be circumvented
- Still the **most popular clustering metric** in network analysis

1. Introduction
2. Algorithms
3. Experiments
4. **NetworkKit**
5. Conclusions

```
initialize nodes with unique labels
while labels not stable do
  | parallel for  $v \in V$ 
  | | adopt dominant label in  $N(v)$ 
  | endfor
end
return communities from labels
```



- communities from labelling of node set
- dense subgraphs agree on common label
→ stable distribution emerges
- a local coverage maximizer
 - getting stuck in local optima of coverage is *desired*
→ modularity implicitly maximized
- $O(m)$ time per iteration, few iterations
- purely local updates → high degree of parallelism



[original, sequential algorithm: Raghavan et al. 07]

- adapted to weighted graphs
- **optimizations**
 - active nodes: evaluate v only if labels in $N(v)$ change
 - truncated iterations: stop if only few nodes undecided
- **OpenMP** parallelization
 - better load balancing with `parallel for schedule(guided)` (high-degree nodes)





Figure: Number of active and updated nodes per iteration of **PLP** on a large web graph

- a locally greedy modularity maximizer
 - repeatedly move nodes to neighbor communities
 - coarsen the graph and repeat
- sequential algorithm: well known method for efficiently achieving high modularity values [Blondel et al. 08]
- our parallel design

```
initialize to singletons
move nodes for modularity gain
if communities changed then
    | coarsen graph
    | recursively apply PLM
    | prolong communities
end
return communities
```

- challenge: evaluate and perform node moves in parallel
 - store and update some interim values for Δmod
 - updates need to be protected by locks
- parallel moves may be based on stale values, but self-correction possible



$$\Delta mod(u, C \rightarrow D) = \underbrace{\frac{\omega(u, D \setminus \{u\}) - \omega(u, C \setminus \{u\})}{\omega(E)}}_{\text{coverage}} + \underbrace{\frac{(\text{vol}(C \setminus \{u\}) - \text{vol}(D \setminus \{u\})) \cdot \text{vol}(u)}{2 \cdot \omega(E)^2}}_{\text{expected coverage}}$$


- add a refinement phase on every level
 - → additional opportunities for modularity improvement at the cost of more iterations

```
initialize to singletons
move nodes for modularity gain
if communities changed then
    | coarsen graph
    | recursively apply PLMR
    | prolong communities
    | move nodes for modularity gain
end
return communities
```


- **ensemble learning**: combine multiple weak classifiers to form a strong one
- generic scheme with exchangeable base and final algorithm
[e.g. Ovelgönne & Geyer-Schulz 13]
 1. ensemble of base algorithms operate on input graph independently
 2. consensus solution is formed and graph coarsened accordingly
 3. final algorithm operates on coarsened graph


```
parallel for Base in ensemble
|    $\zeta_i \leftarrow \mathbf{Base}_i(G)$ 
endfor
 $\bar{\zeta} \leftarrow \mathbf{consensus}(\zeta_1, \dots, \zeta_b)$ 
 $G^1 \leftarrow \mathbf{coarsen}(G, \bar{\zeta})$ 
 $\zeta^1 \leftarrow \mathbf{Final}(G^1)$ 
 $\zeta \leftarrow \mathbf{prolong}(\zeta^1, G)$ 
return  $\zeta$ 
```

- nested parallelism in the ensemble
- efficiently calculate consensus communities through ***k*-way hashing** of community IDs
- base algorithm: focus on speed
- final algorithm: focus on quality optimization




consensus communities [image: Ovelgönne & Geyer-Schulz 13](#)

- nested parallelism in the ensemble
- efficiently calculate consensus communities through ***k*-way hashing** of community IDs
- base algorithm: focus on speed → **PLP**
- final algorithm: focus on quality optimization → **PLM**



consensus communities [image: Ovelgönne & Geyer-Schulz 13](#)



1. Introduction
2. Algorithms
3. Experiments
4. NetworkKit
5. Conclusions

- variety of real-world and synthetic data sets
- **complex networks**: web graphs, internet topology, online social network, scientific collaboration, ...
- **Stochastic Kronecker Graphs (SKG)** for scaling experiments

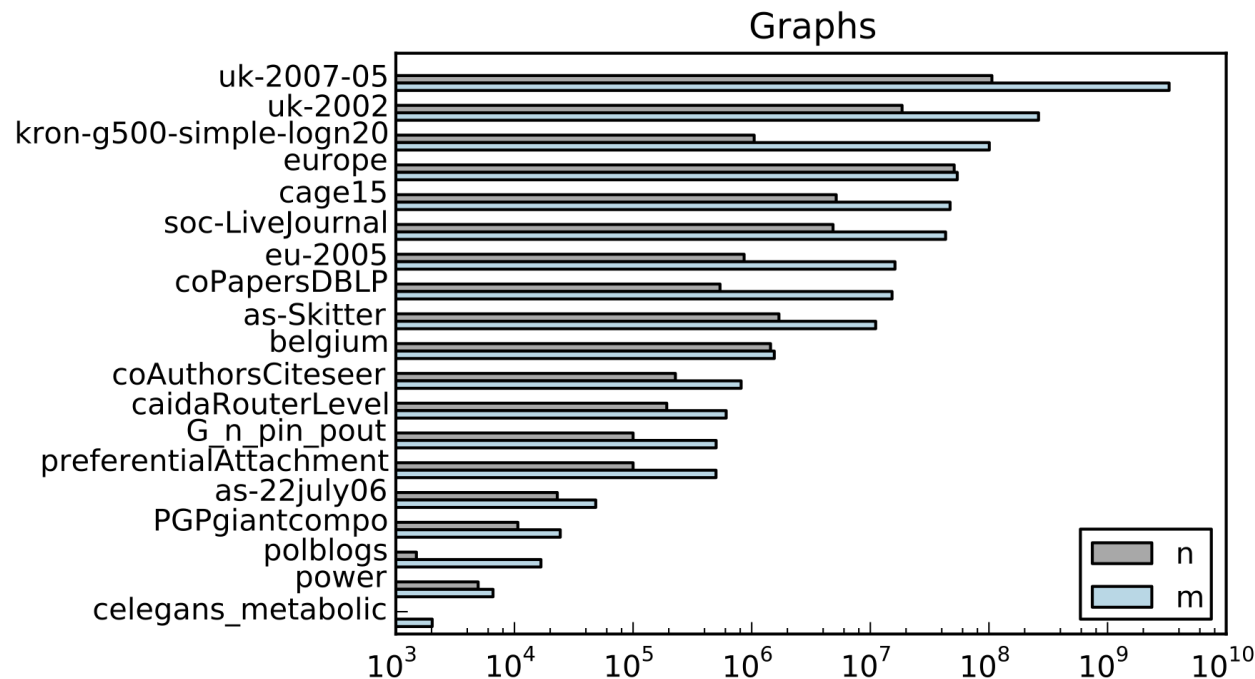


Figure: size comparison of test graphs

	phipute1.itl.kit.edu
compiler	gcc 4.7.1
CPU	2 x 8 Cores: Intel(R) Xeon(R) E5-2680 0 @ 2.70GHz, 32 threads
RAM	256 GB
OS	SUSE 12.2-64

- handles large graphs easily
 - 3.3 billion edge web graph in 60 s with 32 threads
- reasonable modularity values (but room for improvement)




Figure: running time [s] for various networks

- handles large graphs easily
 - 3.3 billion edge web graph in 60 s with 32 threads
- reasonable modularity values (but room for improvement)





Figure: running time [s] for various networks



- Considering the complex input, **PLP** scales well from 2 to 32 threads

Figure: strong scaling of **PLP** on 250M edge web graph

- only minor differences in solution quality between sequential and parallel versions
 - **PLM** able to correct undesirable decisions due to stale data
- better modularity than **PLP** (ca. 0.1)
- but slower (ca. factor 10)
- scaling: worse than **PLP** mainly because of sequential coarsening

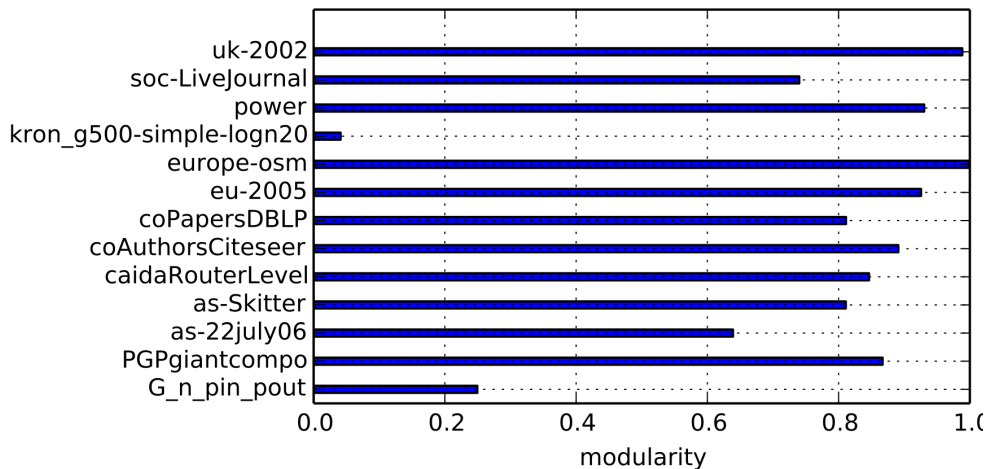


Figure: modularity for **PLM**

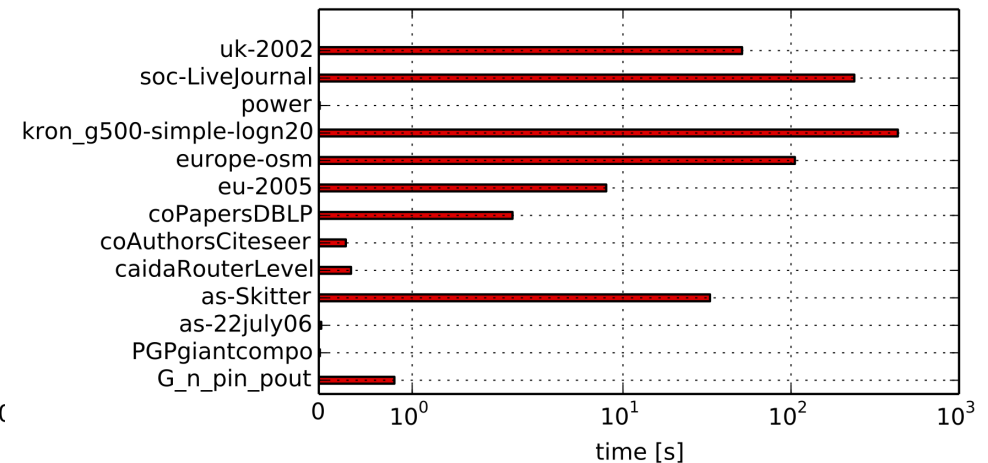


Figure: time for **PLM**

- refinement phase gives small quality boost at the cost of a few more iterations

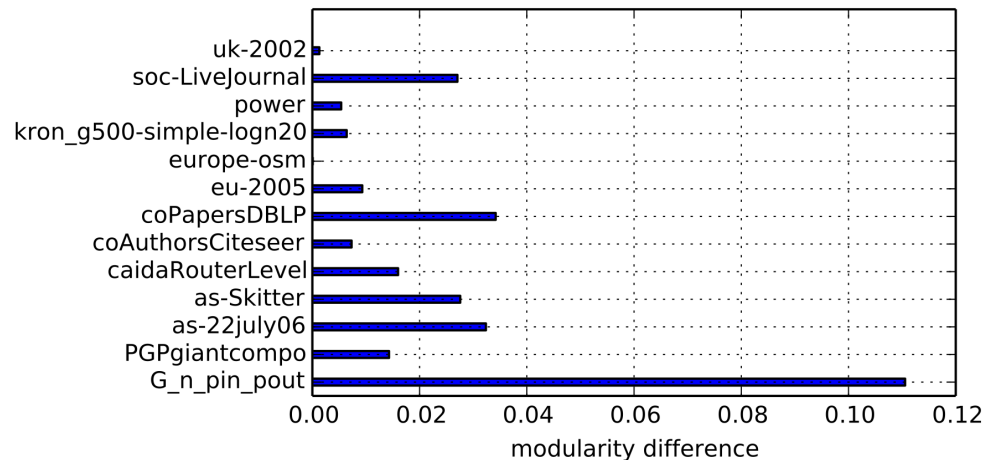


Figure: modularity improvement of **PLMR** compared to **PLM**

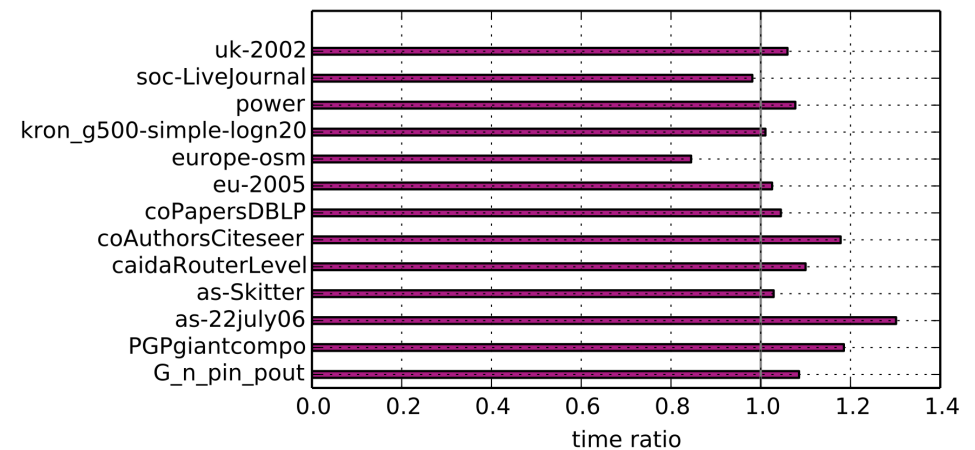


Figure: relative difference in running time of **PLMR** compared to **PLM**

- improved solution quality compared to **PLP**
- small ensembles work best (here: 4-piece ensemble)
- ca. factor 10 slower than **PLP** alone
- modularity improvement ca. 0.05
- slightly faster (and with smaller memory footprint) than **PLM** and **PLMR**

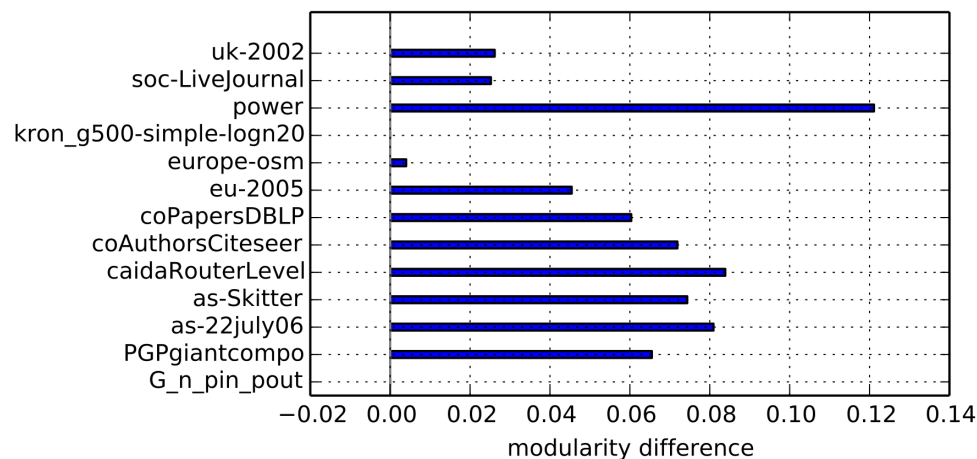


Figure: modularity improvement of **EPP** compared to single **PLP**

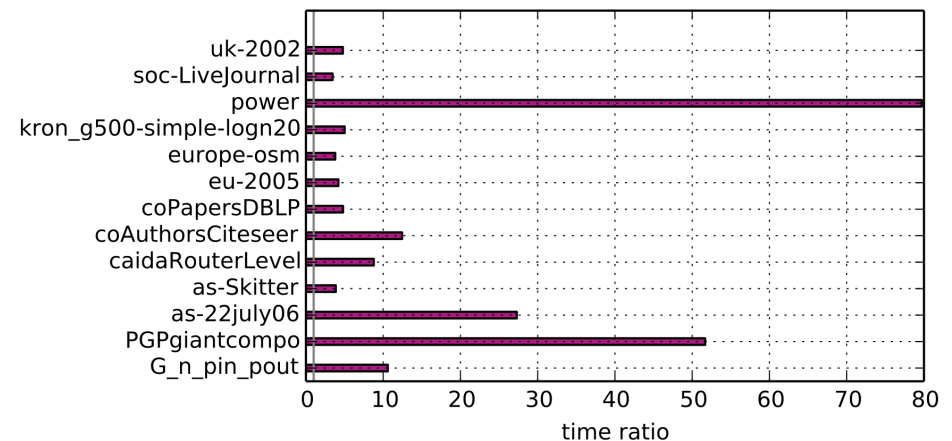



Figure: relative running time of **EPP** compared to single **PLP**

- modularity score: arithmetic mean over all networks of modularity differences
- time score: geometric mean over all networks of relative time differences
- baseline: **PLM** at (0, 1)



1. Introduction
2. Algorithms
3. Experiments
4. **NetworKit**
5. Conclusions


- a toolkit of **high-performance network analysis algorithms**
 - high-performance kernel in **C++11** and **OpenMP**
 - **Python** shell for interactive data analysis (via **Cython**)
- **free software** (*MIT License*)
 - 1.0 (spring 2013): community detection algorithms, data structures
 - 2.0 (November 2013): interactive Python shell
 - 2.1 (TBA): adds various network analysis kernels

Read Network from File

```
In [4]: G = readGraph("pgp.graph")
```

```
In [12]: xscale("log")
         yscale("log")
         plot(properties.degreeDistribution(G))
```

```
Out[12]: [<matplotlib.lines.Line2D at 0x107c5e090>]
```



Network Properties Overview

```
In [5]: properties.showProperties(G)
```

Network Properties

=====

Basic Properties

```
-----
nodes (n)      10680
edges (m)      24316
min. degree    1
max. degree    205
avg. degree    4.55356
isolated nodes 0
self-loops     0
density        0.000426
-----
```

Path Structure

```
-----
connected components      1
size of largest component 10680
diameter
avg. eccentricity
-----
```

Miscellaneous

```
-----
degree assortativity      0.238211
cliques                   13814
-----
```

Community Structure

```
-----
avg. local clustering coefficient      0.265945
PLP community detection
communities      985
modularity       0.795751
PLM community detection
communities      108
modularity       0.879849
-----
```

Community Detection

```
In [13]: communities = community.detectCommunities(G, algo=community.PLM())
```


```
PLM(balanced) detected communities in 0.1455368995666504 [s]  
solution properties:
```

# communities	98
min community size	6
max community size	683
avg. community size	108.98
imbalance	6.26606
modularity	0.882828


- growing collection of network analysis kernels, graph generators, basic graph algorithms etc.
- integration with Python tools for data analysis and visualization
- users and contributors welcome

```
In [7]: communitySizeDistribution = communities.clusterSizes()  
communitySizeDistribution.sort()  
plot(communitySizeDistribution)
```

```
Out[7]: [<matplotlib.lines.Line2D at 0x10c2bfd0>]
```



```
In [8]: viztasks.drawCommunityGraph(G, communities)
```



<http://parco.it.kit.edu/software/networkit.shtml>

1. Introduction
2. Algorithms
3. Experiments
4. **NetworkKit**
5. **Conclusions**

Summary

- developed, implemented and evaluated scalable heuristics for community detection
 - **PLP** extremely fast, but quality may not be high enough
 - **PLM** yields high quality
 - **PLMR** increases quality at the expense of more iterations
 - **EPP** combines their strengths

Ongoing and Future Work

- improve global community detection methods
 - e.g. parallel coarsening
- algorithms for related scenarios
 - selective and dynamic community detection
[presented at ECDA2013 Luxembourg]

Thank you for your attention

Further Reading

C.L. Staudt, H. Meyerhenke:

Engineering High-Performance Community Detection Heuristics for Massive Graphs.
In Proc. 42nd International Conference on Parallel Processing (ICPP 2013).

Acknowledgements

This work was partially funded through the project *Parallel Analysis of Dynamic Networks - Algorithm Engineering of Efficient Combinatorial and Numerical Methods* by the *Ministry of Science, Research and Arts Baden-Württemberg*